


# Chapter Overview

## Games

- 
- A painting of five dogs sitting around a poker table. The dogs are of various breeds, including Beagles and a German Shepherd. They are holding cards and chips, and one dog is dealing. The background features a painting of a ship on the wall and a clock on the right.
- ❖ Motivation
  - ❖ Objectives
  - ❖ Games and AI
  - ❖ Games and Search
  - ❖ Perfect Decisions
  - ❖ Imperfect Decisions
  - ❖ Alpha-Beta Pruning
  - ❖ Games with Chance
  - ❖ Games and Computers
  - ❖ Important Concepts and Terms
  - ❖ Chapter Summary



# Motivation

- ❖ **examine the role of AI methods in games**
- ❖ **some game provide challenges that can be formulated as abstract competitions with clearly defined states and rules**
  - ❖ programs for some games can be derived from search methods
  - ❖ narrow view of games
- ❖ **games can be used to demonstrate the power of computer-based techniques and methods**
- ❖ **more challenging games require the incorporation of specific knowledge and information**
- ❖ **expansion of the use of games**
  - ❖ from entertainment to training and education

# Objectives

- ❖ **explore the combination of AI and games**
- ❖ **understand the use and application of search methods to game programs**
  - ❖ apply refined search methods such as minimax to simple game configurations
  - ❖ use alpha-beta pruning to improve the efficiency of game programs
  - ❖ understand the influence of chance on the solvability of chance-based games
- ❖ **evaluation of methods**
  - ❖ suitability of game techniques for specific games
  - ❖ suitability of AI methods for games

# Games and Computers

- ❖ **games offer concrete or abstract competitions**
  - ❖ “I’m better than you!”
- ❖ **some games are amenable to computer treatment**
  - ❖ mostly mental activities
  - ❖ well-formulated rules and operators
  - ❖ accessible state
- ❖ **others are not**
  - ❖ emphasis on physical activities
  - ❖ rules and operators open to interpretation
    - ❖ need for referees, mitigation procedures
  - ❖ state not (easily or fully) accessible

# Games and AI

- ❖ **traditionally, the emphasis has been on a narrow view of games**
  - ❖ formal treatment, often as an expansion of search algorithms
- ❖ **more recently, AI techniques have become more important in computer games**
  - ❖ computer-controlled characters (agents)
  - ❖ adaptive performance
  - ❖ more sophisticated story lines
  - ❖ more complex environments
  - ❖ better overall user experience
- ❖ **books exist about Games & AI**
  - ❖ collections at
    - ❖ [The AI Programmers Bookshelf](#)
    - ❖ [Recommended Ai Books And Sites](#) by Dave Mark at [gamedev.net](http://gamedev.net)
    - ❖ [386 Game AI Programming Articles and Counting...](#)

# Cognitive Game Design

## ❖ **story development**

- ❖ generation of interesting and appealing story lines
- ❖ variations in story lines
- ❖ analysis of large-scale game play

## ❖ **character development**

- ❖ modeling and simulation of computer-controlled agents
- ❖ possibly enhancement of user-controlled agents

## ❖ **immersion**

- ❖ strong engagement of the player's mind

## ❖ **emotion**

- ❖ integration of plausible and believable motion in characters
- ❖ consideration of the user's emotion

## ❖ **pedagogy**

- ❖ achievement of “higher” goals through entertainment

# Game Analysis

- ❖ **often deterministic**

- ❖ the outcome of actions is known
- ❖ sometimes an element of chance is part of the game
  - ❖ e.g. dice

- ❖ **two-player, turn-taking**

- ❖ one move for each player

- ❖ **zero-sum utility function**

- ❖ what one player wins, the other must lose

- ❖ **often perfect information**

- ❖ fully observable, everything is known to both players about the state of the environment (game)
- ❖ not for all games
  - ❖ e.g. card games with “private” or “hidden” cards
  - ❖ Scrabble

# Games as Adversarial Search


- ❖ **many games can be formulated as search problems**
- ❖ **the zero-sum utility function leads to an adversarial situation**
  - ❖ in order for one agent to win, the other necessarily has to lose
- ❖ **factors complicating the search task**
  - ❖ potentially huge search spaces
  - ❖ elements of chance
  - ❖ multi-person games, teams
  - ❖ time limits
  - ❖ imprecise rules



# The Ideal of Non-Adversarial Games

## ❖ Lab 10 Submission: AI and Humor -> Game Theory AI

❖ by Connor Citron - Monday, November 19, 2012, 7:45 PM



# Difficulties with Games

- ❖ **games can be very hard search problems**

- ❖ yet reasonably easy to formalize
- ❖ finding the optimal solution may be impractical
  - ❖ a solution that beats the opponent is “good enough”
- ❖ unforgiving
  - ❖ a solution that is “not good enough” leads to higher costs, and to a loss to the opponent

- ❖ **example: chess**

- ❖ size of the search space
  - ❖ branching factor around 35
  - ❖ about 50 moves per player
  - ❖ about  $35^{100}$  or  $10^{154}$  nodes
    - ❖ about  $10^{40}$  distinct nodes (size of the search graph)

# Games and Search

- ❖ the actions of an agent playing a game can often be formulated as a search problem
- ❖ some factors make the use of search methods challenging
  - ❖ multiple players
  - ❖ contingencies
    - ❖ non-deterministic events
    - ❖ actions of opponents
    - ❖ chance events (e.g. dice)
  - ❖ consideration of probabilities
  - ❖ ...

# Search Problem Formulation

- ❖ **initial state**
  - ❖ board, positions of pieces
  - ❖ whose turn is it
- ❖ **successor function (operators)**
  - ❖ list of (move, state)
  - ❖ defines the legal moves, and the resulting states
    - ❖ determines the rules of the game
- ❖ **terminal test**
  - ❖ also called goal test
  - ❖ determines when the game is over (terminal state)
  - ❖ calculates the result
    - ❖ usually win, lose, draw; sometimes a score (utility function)
- ❖ **utility or payoff function**
  - ❖ numeric value for the outcome of a game



# Utility Function

- ❖ **also called objective function, payoff function**
- ❖ **final numeric value at a terminal state**
  - ❖ should be consistent with the “outcome” of the game
    - ❖ win, lose, tie
    - ❖ score
- ❖ **zero-sum games**
  - ❖ the sum of the utility values for each player is the same for every instance of the game

# Evaluation Function

- ❖ **an estimate of the expected utility value for a node in the game state**
- ❖ **used instead of the utility function when it is impractical to calculate it**
  - ❖ should be consistent with the utility function
    - ❖ same ordering of the terminal nodes
  - ❖ should reflect the chances of winning
- ❖ **usually based on heuristics**
  - ❖ knowledge about the domain (game)
  - ❖ previous experience
- ❖ **often relies on “features”**
  - ❖ important aspects of the game
    - ❖ number of pieces, positions on the board, time constraints
  - ❖ basis for weighted linear evaluation functions
    - ❖ sum of weighted feature values

# Single-Person Game

- ❖ **conventional search problem**
  - ❖ identify a sequence of moves that leads to a winning state
  - ❖ examples:
    - ❖ Solitaire, dungeons and dragons, Rubik's cube,
  - ❖ little attention in AI
- ❖ **some games can be quite challenging**
  - ❖ some versions of solitaire
  - ❖ Rubik's cube
    - ❖ a heuristic for this was found by the Absolver theorem prover

# Contingency Problem

- ❖ **in general:**
  - ❖ non-deterministic events
  - ❖ unknown outcomes actions
- ❖ **uncertainty due to the moves and motivations of the opponent**
  - ❖ tries to make the game as difficult as possible for the player
    - ❖ attempts to maximize its own utility function value
    - ❖ thus minimize the player's utility function value
  - ❖ different from contingency due to neutral factors, such as
    - ❖ chance
    - ❖ outside influence



# Two-Person Games

- ❖ **games with two opposing players**
  - ❖ often called MIN and MAX
  - ❖ usually MAX moves first, then they take turns
  - ❖ in game terminology, a move comprises two steps (“plies”)
    - ❖ one by MAX and one by MIN
- ❖ **MAX wants a strategy to find a winning state**
  - ❖ no matter what MIN does
- ❖ **MIN does the same**
  - ❖ or at least tries to prevent MAX from winning
- ❖ **full information**
  - ❖ both players know the full state of the environment
- ❖ **partial information**
  - ❖ one player only knows part of the environment
  - ❖ some aspects may be hidden from the opponent, or from both players

# Perfect Decisions

- ❖ **based on an rational (optimal) strategy for MAX**
  - ❖ traverse all relevant parts of the search tree
    - ❖ this must include possible moves by MIN
      - ❖ often MIN is also assumed to act rationally
  - ❖ identify a path that leads MAX to a winning state
  - ❖ based on utility values of nodes
    - ❖ reflects how well the agent is doing at a particular node
- ❖ **often impractical**
  - ❖ time and space limitations
  - ❖ need for a utility function
    - ❖ should reflect the chances of winning
    - ❖ partial information
    - ❖ non-deterministic aspects

# MiniMax Strategy

## ♦ optimal strategy for MAX

❖ not very practical

- generate the whole game tree
- calculate the value of each terminal state
  - based on the utility function
- calculate the utilities of the nodes bottom-up
  - starting from the leaf nodes up to the root
- MAX selects the value with the highest node
- MAX assumes that MIN in its move will select the node that minimizes the value

# MiniMax Value

- ❖ **utility of being in the state that corresponds to a node**
  - ❖ MAX's perspective:
    - ❖ move to a state with the maximum value
      - ❖ highest chance of winning
  - ❖ MIN's perspective:
    - ❖ move to a state with the minimum value
  - ❖ assumes that both players play optimally

```
function MiniMax-Value(state) returns a utility value
  if Terminal-Test(state) then
    return Utility(state)
  else if Max is to move then
    return the highest MiniMax-Value of Successors(state)
  else
    return the lowest MiniMax-Value of Successors(state)
```



# MiniMax Algorithm

- ❖ selects the best successor from a given state
- ❖ invokes MINIMAX-VALUE for each successor state

```
function MiniMax-Decision(state) returns action  
    for each s in Successors[state] do  
        Value[s] := MiniMax-Value(s)  
    end  
    return action with the highest Value[s]
```


# MiniMax Properties

- ❖ **based on depth-first**
  - ❖ recursive implementation
- ❖ **time complexity is  $O(b^m)$** 
  - ❖ exponential in the number of moves
- ❖ **space complexity is  $O(b \cdot m)$** 
  - ❖ inherited from depth-first

b    branching factor


m    maximum depth of the search tree

# MiniMax Example 1



terminal nodes: values calculated from the utility function (made-up, doesn't reflect any real game situation)


# MiniMax Example 2




other nodes: values calculated via minimax algorithm starting at the bottom (terminal nodes)




# MiniMax Example 3




# MiniMax Example 4



# MiniMax Example 5



# MiniMax Example 6



moves by **MAX** and countermoves by **MIN**

# MiniMax Observations

- ❖ the values of some of the leaf nodes are irrelevant for decisions at the next level
- ❖ this also holds for decisions at higher levels
- ❖ as a consequence, under certain circumstances, some parts of the tree can be disregarded
  - ❖ it is possible to still make an optimal decision without considering those parts
  - ❖ these parts can be “pruned”

# Pruning


- ❖ **discards parts of the search tree**
  - ❖ guaranteed not to contain good moves
  - ❖ guarantee that the solution is not in that branch or sub-tree
    - ❖ if both players make optimal (rational) decisions, they will never end up in that part of the search tree
    - ❖ sub-optimal moves by the opponent may lead into that part
      - ❖ may increase the amount of calculations for the player, but does not change the outcome of the game
- ❖ **results in substantial time and space savings**
  - ❖ as a consequence, longer sequences of moves can be explored
  - ❖ the leftover part of the task may still be exponential, however



# Alpha-Beta Pruning

- ❖ **certain moves are not considered**
  - ❖ new states won't result in a better evaluation value than a move further up in the tree
  - ❖ they would lead to a less desirable outcome
- ❖ **applies to moves by both players**
  - ❖  $\alpha$  indicates the best choice for Max so far
    - ❖ never decreases
  - ❖  $\beta$  indicates the best choice for Min so far
    - ❖ never increases
- ❖ **extension of the minimax approach**
  - ❖ results in the same sequence of moves as minimax, but with less overhead
  - ❖ prunes uninteresting parts of the search tree

# Alpha-Beta Example 1




$\alpha$  best choice for **MAX**?

$\beta$  best choice for **MIN** ?


- ◆ we assume a depth-first, left-to-right search as basic strategy
- ◆ the range of the possible values for each node are indicated
  - ❖ initially  $[-\infty, +\infty]$
  - ❖ from **MAX**'s or **MIN**'s perspective
  - ❖ these *local* values reflect the values of the sub-trees in that node; the *global* values  $\alpha$  and  $\beta$  are the best overall choices so far for **MAX** or **MIN**

# Alpha-Beta Example 2



- ◆ **MIN** obtains the first value from a successor node

# Alpha-Beta Example 3




$\alpha$  best choice for **MAX** ?

$\beta$  best choice for **MIN** 6


- ♦ **MIN** obtains the second value from a successor node

# Alpha-Beta Example 4



- ◆ **MIN** obtains the third value from a successor node
  - ◆ this is the last value from this sub-tree, and the exact value is known
- ◆ **MAX** now has a value for its first successor node, but hopes that something better might still come

# Alpha-Beta Example 5




$\alpha$  best choice for **MAX** 5


$\beta$  best choice for **MIN** 3

- ♦ **MIN** continues with the next sub-tree, and gets a better value
- ♦ **MAX** has a better choice from its perspective, however, and will not consider a move in the sub-tree currently explored by **MIN**
  - ❖ initially  $[-\infty, +\infty]$




# Alpha-Beta Example 6



- ♦ **MIN** knows that **MAX** won't consider a move to this sub-tree, and abandons it
- ♦ this is a case of *pruning*, indicated by 

# Alpha-Beta Example 7




$\alpha$  best choice for **MAX**     $5$

$\beta$  best choice for **MIN**     $3$


- ◆ **MIN** explores the next sub-tree, and finds a value that is worse than the other nodes at this level
- ◆ if **MIN** is not able to find something lower, then **MAX** will choose this branch, so **MIN** must explore more successor nodes

# Alpha-Beta Example 8



- ♦ **MIN** is lucky, and finds a value that is the same as the current worst value at this level
- ♦ **MAX** can choose this branch, or the other branch with the same value

# Alpha-Beta Example 9



$\alpha$  best choice for **MAX** 5

$\beta$  best choice for **MIN** 3

- ◆ **MIN** could continue searching this sub-tree to see if there is a value that is less than the current worst alternative in order to give **MAX** as few choices as possible
  - ❖ this depends on the specific implementation
- ◆ **MAX** knows the best value for its sub-tree

# Alpha-Beta Algorithm

```
function Max-Value(state, alpha, beta) returns a utility value
  if Terminal-Test (state) then return Utility(state)
  for each s in Successors(state) do
    alpha := Max (alpha, Min-Value(s, alpha, beta))
    if alpha >= beta then return beta
  end
  return alpha
```

```
function Min-Value(state, alpha, beta) returns a utility value
  if Terminal-Test (state) then return Utility(state)
  for each s in Successors(state) do
    beta := Min (beta, Max-Value(s, alpha, beta))
    if beta <= alpha then return alpha
  end
  return beta
```



# Properties of Alpha-Beta Pruning

- ❖ **in the ideal case, the best successor node is examined first**
  - ❖ results in  $O(b^{d/2})$  nodes to be searched instead of  $O(b^d)$ 
    - ❖ alpha-beta can look ahead “twice as far” as minimax
  - ❖ in practice, simple ordering functions are quite useful
- ❖ **assumes an idealized tree model**
  - ❖ uniform branching factor, path length
  - ❖ random distribution of leaf evaluation values
- ❖ **transpositions tables can be used to store permutations**
  - ❖ sequences of moves that lead to the same position
- ❖ **requires additional information for good players**
  - ❖ game-specific background knowledge
  - ❖ empirical data



# Imperfect Decisions

- ❖ **complete search is impractical for most games**
- ❖ **alternative: search the tree only to a certain depth**
  - ❖ requires a cutoff-test to determine where to stop
    - ❖ replaces the terminal test
    - ❖ the nodes at that level effectively become terminal leaf nodes
  - ❖ uses a heuristics-based evaluation function to estimate the expected utility of the game from those leaf nodes


# Evaluation Function

- ❖ **determines the performance of a game-playing program**
- ❖ **must be consistent with the utility function**
  - ❖ values for terminal nodes (or at least their order) must be the same
- ❖ **tradeoff between accuracy and time cost**
  - ❖ without time limits, minimax could be used
- ❖ **should reflect the actual chances of winning**
- ❖ **frequently weighted linear functions are used**
  - ❖  $E = w_1 f_1 + w_2 f_2 + \dots + w_n f_n$
  - ❖ combination of features, weighted by their relevance


# Example: Tic-Tac-Toe

- ❖ **simple evaluation function**
- ❖  **$E(s) = (rx + cx + dx) - (ro + co + do)$** 
  - ❖ where r,c,d are the numbers of row, column and diagonal lines still available;  
x and o are the pieces of the two players
- ❖ **1-ply lookahead**
  - ❖ start at the top of the tree
  - ❖ evaluate all 9 choices for player 1
  - ❖ pick the maximum E-value
- ❖ **2-ply lookahead**
  - ❖ also looks at the opponent's possible move
    - ❖ assumes rational behavior: the opponent picks the minimum E-value

# Tic-Tac-Toe 1-Ply



# Tic-Tac-Toe 2-Ply



# Checkers Case Study

## ❖ initial board configuration

❖ **BLACK**      single on 20  
                      single on 21  
                      king on 31

❖ **RED**            single on 23  
                          king on 22

evaluation function

$$E(s) = (5x_1 + x_2) - (5r_1 + r_2)$$


where

$x_1$  = black king advantage,

$x_2$  = black single advantage,



$r_1$  = red king advantage,

$r_2$  = red single advantage







# Example




# Checkers Alpha-Beta Example




# Checkers Alpha-Beta Example



⊘  **$\beta$ -cutoff**: no need to examine further branches



# Checkers Alpha-Beta Example





# Checkers Alpha-Beta Example



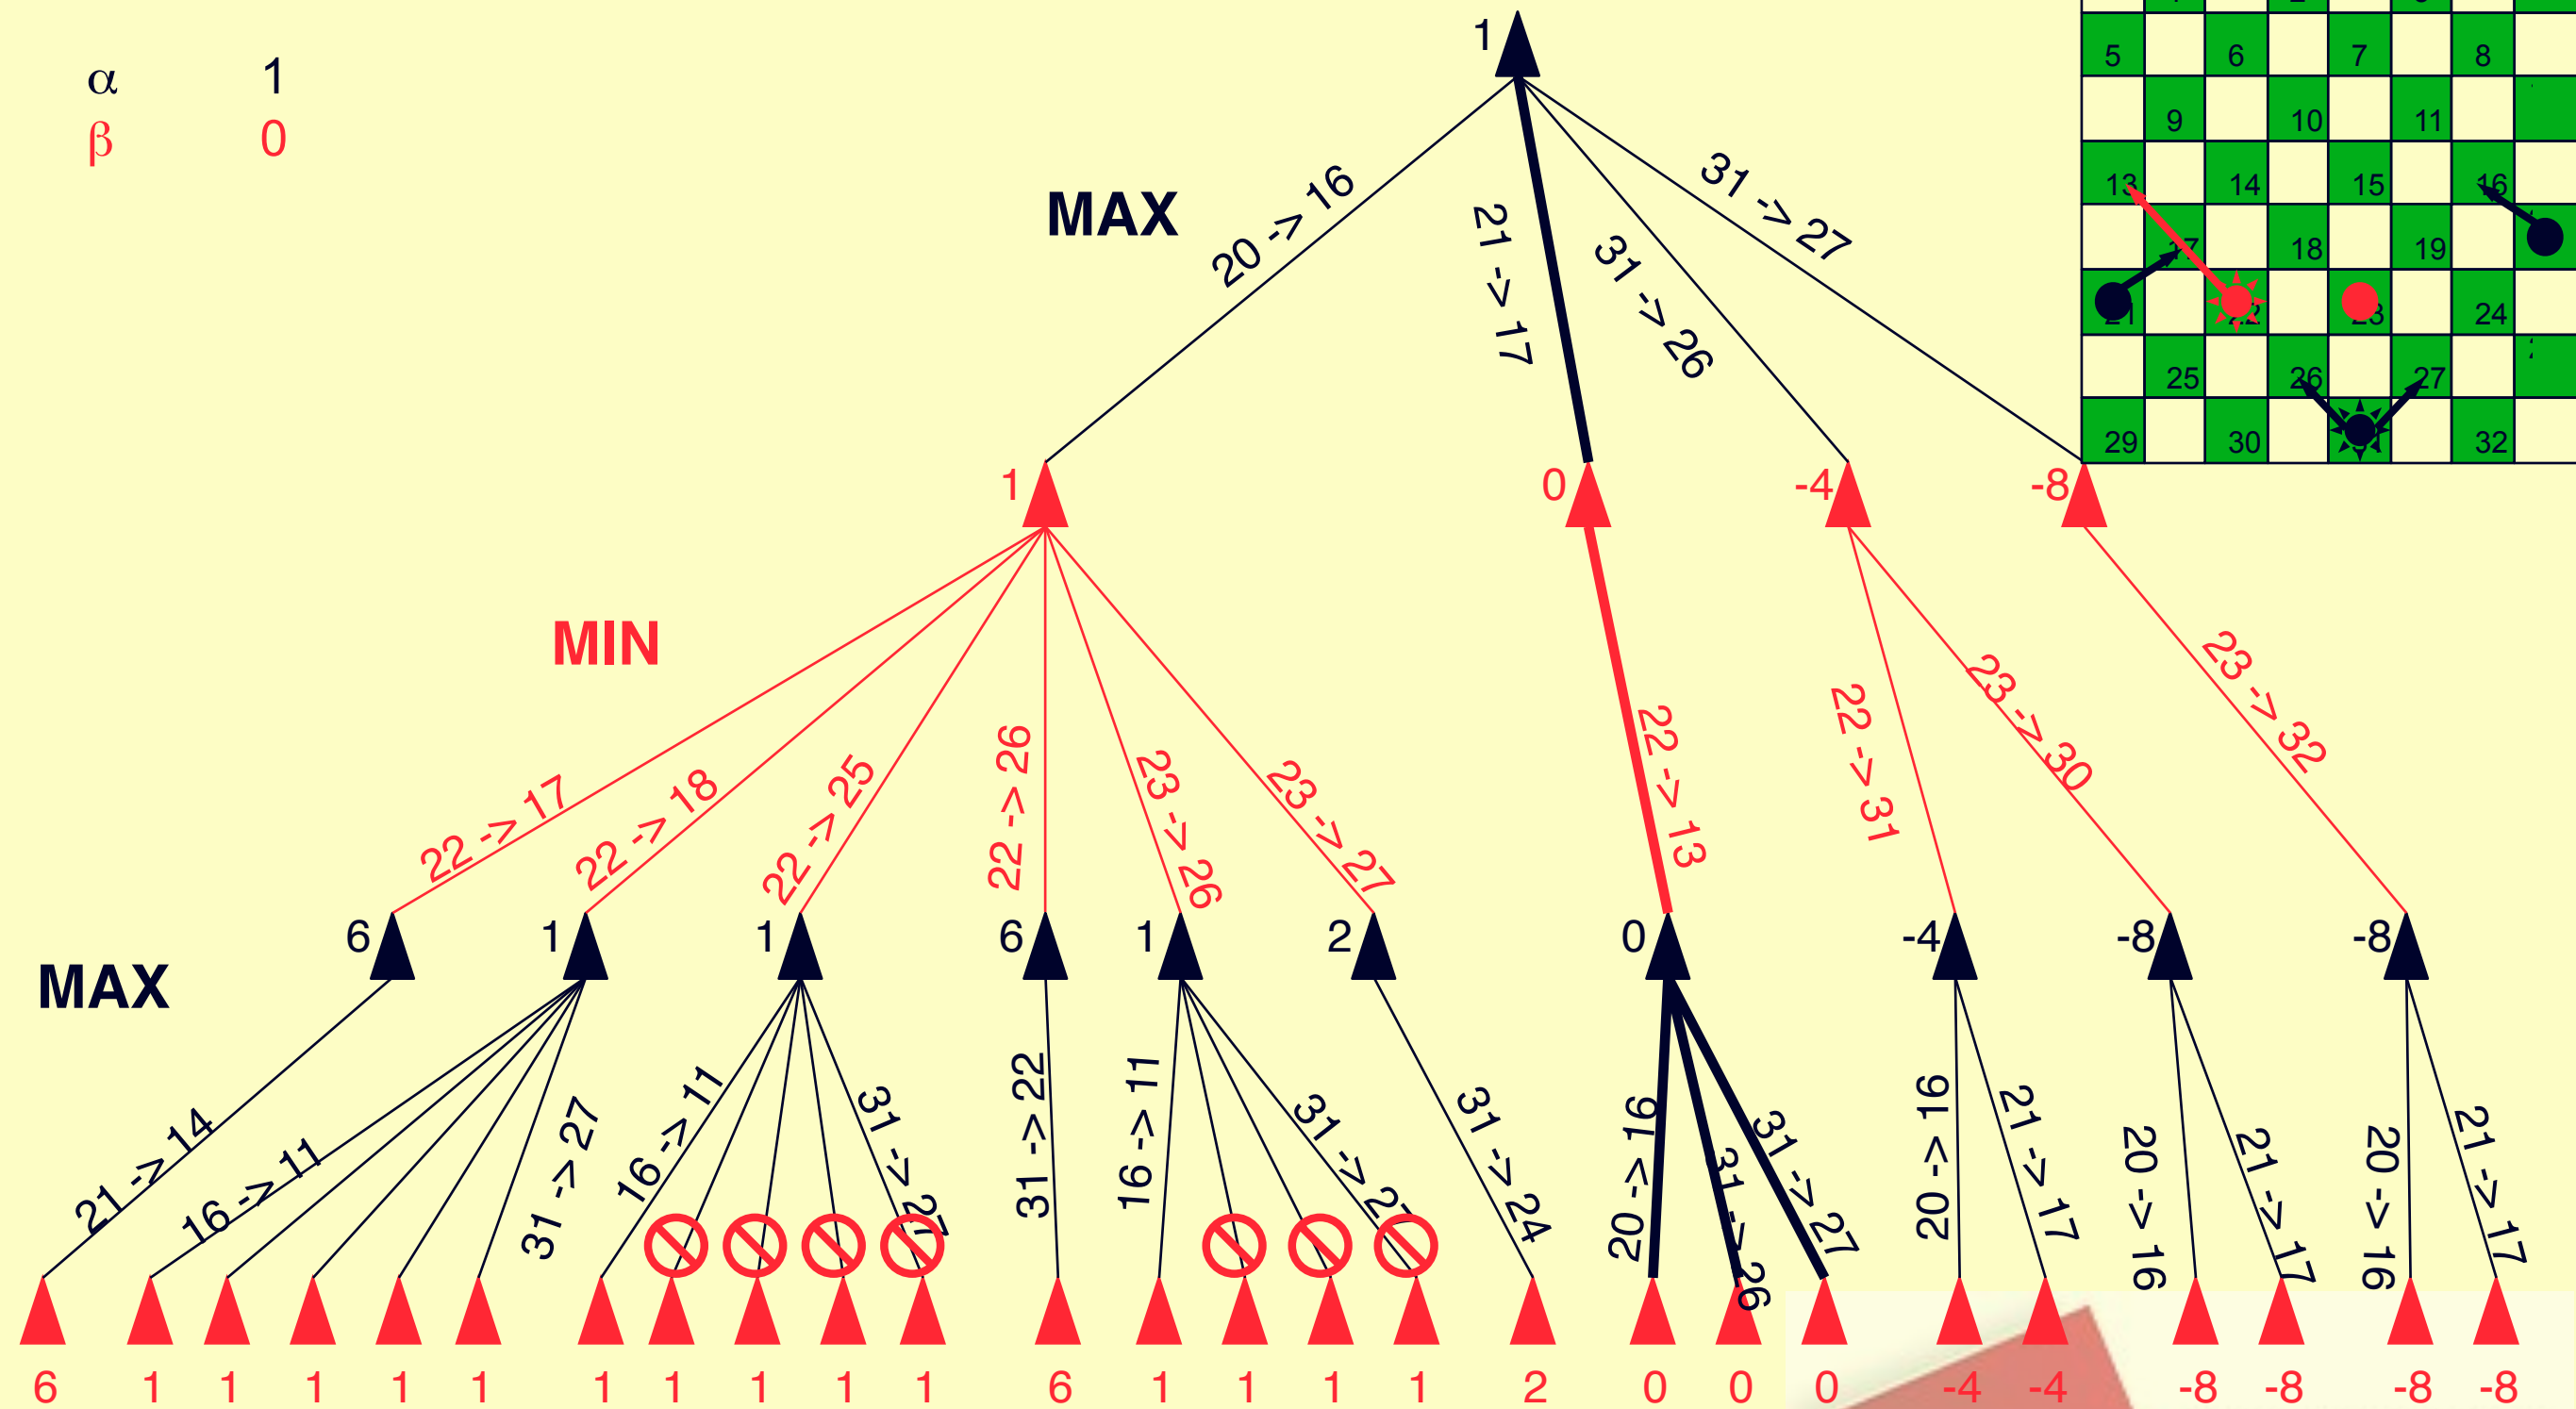





Diagram 1: An 8x8 grid with alternating yellow and green squares. Numbers 1 through 32 are placed in the yellow squares in row-major order. A red sun icon is on square 22, a red circle icon is on square 23, and a blue star icon is on square 32. Arrows indicate movement: a black arrow from 16 to 19, a red arrow from 23 to 27, and a black arrow from 32 to 27.




# Checkers Alpha-Beta Example



# Checkers Alpha-Beta Example



# Checkers Alpha-Beta Example



# Search Limits

- ❖ **search must be cut off because of time or space limitations**
- ❖ **strategies like depth-limited or iterative deepening search can be used**
  - ❖ don't take advantage of knowledge about the problem
- ❖ **more refined strategies apply background knowledge**
  - ❖ quiescent search
    - ❖ cut off only parts of the search space that don't exhibit big changes in the evaluation function

# Horizon Problem

- ❖ **moves may have disastrous consequences in the future, but the consequences are not visible**
  - ❖ the corresponding change in the evaluation function will only become evident at deeper levels
    - ❖ they are “beyond the horizon”
- ❖ **determining the horizon is an open problem without a general solution**
  - ❖ only some pragmatic approaches restricted to specific games or situation



# Games with Chance

- ❖ **in many games, there is a degree of unpredictability through random elements**
  - ❖ throwing dice, card distribution, roulette wheel, ...
- ❖ **this requires chance nodes in addition to the Max and Min nodes**
  - ❖ branches indicate possible variations
  - ❖ each branch indicates the outcome and its likelihood

# Rolling Dice

- ❖ **36 ways to roll two dice**
  - ❖ the same likelihood for all of them
  - ❖ due to symmetry, there are only 21 distinct rolls
  - ❖ six doubles have a  $1/36$  chance
  - ❖ the other fifteen have a  $1/18$  chance

# Decisions with Chance

- ❖ **the utility value of a position depends on the random element**
  - ❖ the definite minimax value must be replaced by an expected value
- ❖ **calculation of expected values**
  - ❖ utility function for terminal nodes
  - ❖ for all other nodes
    - ❖ calculate the utility for each chance event
    - ❖ weigh by the chance that the event occurs
    - ❖ add up the individual utilities

# Expectiminimax Algorithm

- ❖ calculates the utility function for a particular position based on the outcome of chance events
  - ❖ utilizes an additional pair of functions to assess the utility values of chance nodes
  - ❖  $\text{expectimin}(C) = \sum_i P(d_i) \min_{s \in S(C, d_i)} (\text{utility}(s))$
  - ❖  $\text{expectimax}(C) = \sum_i P(d_i) \max_{s \in S(C, d_i)} (\text{utility}(s))$
- ❖ where  $C$  are chance nodes,
- ❖  $P(d_i)$  is the probability of a chance event  $d_i$ , and  $S(C, d_i)$  the set of positions resulting from the event  $d_i$ , occurring at position  $C$

# Limiting Search with Chance

- ❖ **similar to alpha-beta pruning for minimax**
  - ❖ search is cut off
  - ❖ evaluation function is used to estimate the value of a position
  - ❖ must put boundaries on possible values of the utility function
- ❖ **somewhat more restricted**
  - ❖ the evaluation function is influenced by some aspects of the chance events

# Properties of Expectiminimax

- ❖ **complexity of  $O(b^m n^m)$** 
  - ❖  $n$  - number of distinct chance events
  - ❖  $b$  - branching factor
  - ❖  $m$  - maximum path length (number of moves in the game)
- ❖ example backgammon:
  - ❖  $n = 21$ ,  $b \approx 20$  (but may be as high as 4000)



# Games and Computers

## ❖ state of the art for some game programs

- ❖ Chess
- ❖ Checkers
- ❖ Othello
- ❖ Backgammon
- ❖ Go

# Chess

- ❖ **Deep Blue, a special-purpose parallel computer, defeated the world champion Gary Kasparov in 1997**
  - ❖ the human player didn't show his best game
    - ❖ some claims that the circumstances were questionable
  - ❖ Deep Blue used a massive data base with games from the literature
- ❖ **Fritz, a program running on an ordinary PC, challenged the world champion Vladimir Kramnik to an eight-game draw in 2002**
  - ❖ top programs and top human players are roughly equal
- ❖ **Houdini**
  - ❖ for a development of the strongest chess engine ~2012, see <http://www.chessbase.com/newsdetail.asp?newsid=8591>
- ❖ **best players become “hybrid”**
  - ❖ human supported by computers

# Checkers

- ❖ **Arthur Samuel developed a checkers program in the 1950s that learns its own evaluation function**
  - ❖ reached an expert level stage in the 1960s
- ❖ **Chinook became world champion in 1994**
  - ❖ human opponent, Dr. Marion Tinsley, withdrew for health reasons
    - ❖ Tinsley had been the world champion for 40 years
    - ❖ he lost only seven games (two of them to the Chinook computer program) in his entire 45-year career
  - ❖ Chinook used off-the-shelf hardware, alpha-beta search, end-games data base for six-piece positions

# Othello

- ❖ **Logistello defeated the human world champion in 1997**
- ❖ **many programs play far better than humans**
  - ❖ smaller search space than chess
  - ❖ little evaluation expertise available

# Backgammon

- ❖ **TD-Gammon, a neural-network based program developed in 1992, ranked among the best players in the world**
  - ❖ improved its own evaluation function through temporal difference learning techniques by playing against clones of itself
  - ❖ two-ply lookahead search
    - ❖ drawback is poor end play
  - ❖ started out “knowledge-free”
    - ❖ no experience through prior games
  - ❖ its success with unorthodox strategies had a significant impact on the backgammon community
    - ❖ since it examined moves previously not considered by players
  - ❖ some improvement in a later version that included expert-designed features
  - ❖ search-based methods are practically hopeless for backgammon
    - ❖ chance elements, branching factor

# Go

- ❖ **humans play far better**
  - ❖ large branching factor (around 360)
    - ❖ search-based methods are hopeless
- ❖ **rule-based systems play at amateur level**
- ❖ **the use of pattern-matching techniques can improve the capabilities of programs**
  - ❖ difficult to integrate
- ❖ **\$2,000,000 prize for the first program to defeat a top-level player**

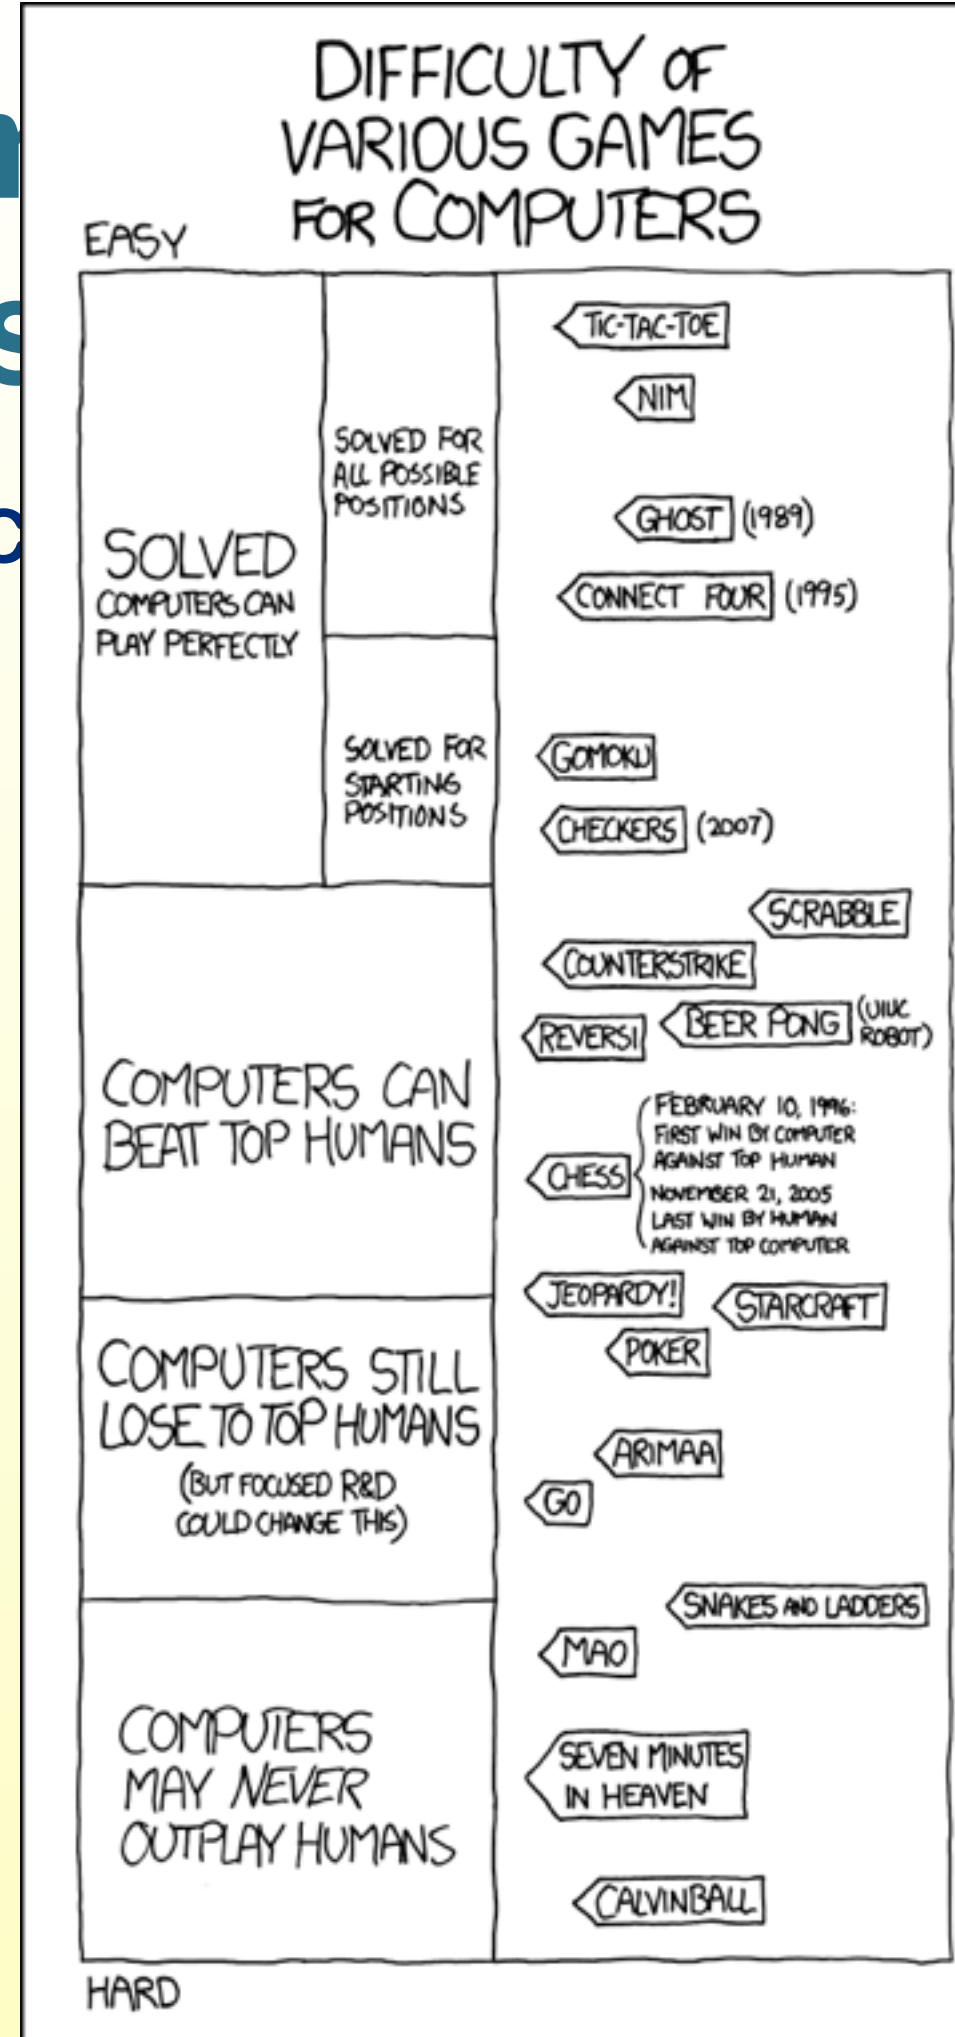
# Jeopardy

- ❖ in 2010, IBM announced that its Watson system will participate in a Jeopardy contest
- ❖ Watson beat two of the best Jeopardy participants



# Difficulty of Games for Computers

- ❖ Lab 10 Submission: AI and Humor -> XKCD
- ❖ by Andrew Guenther - Tuesday, November 20,



# Beyond Search?

- ❖ **search-based game playing strategies have some inherent limitations**
  - ❖ high computational overhead
  - ❖ exploration of uninteresting areas of the search space
  - ❖ complicated heuristics
- ❖ **utility of node expansion**
  - ❖ consider the trade-off between the costs for calculations, and the improvement in traversing the search space
- ❖ **goal-based reasoning and planning**
  - ❖ concentrate on possibly distant, but critical states instead of complete paths with lots of intermediate states
- ❖ **meta-reasoning**
  - ❖ observe the reasoning process itself, and try to improve it
  - ❖ alpha-beta pruning is a simple instance

# Important Concepts and Terms

- ❖ action
- ❖ alpha-beta pruning
- ❖ Backgammon
- ❖ chance node
- ❖ Checkers
- ❖ Chess
- ❖ contingency problem
- ❖ evaluation function
- ❖ expectiminimax algorithm
- ❖ Go
- ❖ heuristic
- ❖ horizon problem
- ❖ initial state
- ❖ minimax algorithm
- ❖ move
- ❖ operator
- ❖ Othello
- ❖ ply
- ❖ pruning
- ❖ quiescent
- ❖ search
- ❖ search tree
- ❖ state
- ❖ strategy
- ❖ successor
- ❖ terminal state
- ❖ utility function

# Chapter Summary

- ❖ many game techniques are derived from search methods
- ❖ the minimax algorithm determines the best move for a player by calculating the complete game tree
- ❖ alpha-beta pruning dismisses parts of the search tree that are provably irrelevant
- ❖ an evaluation function gives an estimate of the utility of a state when a complete search is impractical
- ❖ chance events can be incorporated into the minimax algorithm by considering the weighted probabilities of chance events

